

Problem Statement

Problem Statement may undergo changes as it still evolving. Please be alert to updates and check back frequently.

Write a simulator for a multiple server queuing system, with servers that change speed dynamically (for power saving).

Inputs:

- Number of servers
- Buffer size (maximum number of requests that are waiting in the system not including those in service). Requests arriving when buffer is full are dropped.
- Min speed, Max speed that the server can operate at (in MIPS - instructions per second)
- step size S (in mips)
- probe interval P (in seconds). The job queue size is checked every P seconds and speed change action may be taken (see below)
- Job Queue size thresholds: Jlow, JHigh. At a probe, whenever job queue length (waiting requests only) is seen above threshold JHigh, speed is increased by S (upto max). Whenever job queue length is seen below threshold Jlow, speed is decreased by S (upto min).
- A "trace" of request arrivals with the following information:
request_arrival_time request_processing_size_in_million_instructions

Trace Outputs:

- A request trace file (request.log) with the following format:
Request _# arrival_time processing_start_time processing_end_time waiting_time service_time
- A server trace file (server.log) with the following format
probe_interval_# job_q_size_at_probe speed_before_probe speed_after_probe
- An event trace written to the screen: where a line of output is generated just after an event has been removed from the event list as the next event to be processed. At this point, the following line should be written out:
sim_time event_being_processed next_req_in_buffer total_num_waiting_in_buffer next_event_type

Here, event_being_processed will be : ARRIVAL/DEPARTURE/PROBE

next_req_in_buffer should be shown as: (A1,S1) where (A1, S1) is the first request in the buffer queue with arrival time A1 and service requirement of S1.

next_event_type will also be ARRIVAL/DEPARTURE/PROBE

Summarized outputs (write to screen):

Number of requests arrived (including dropped):

Number of requests serviced:

Number of requests dropped:

Average waiting time:
Average service time:
90%ile of waiting time:
Maximum waiting time:
Max queue length reached:
Average operating speed of the server:
90%ile of speed:

Bonus Points for:

Simulating request timeouts (sample input for this TBD)
Average Server Utilization
Average Job queue length

[Sample Input](#)

[Sample Output](#) (Format only, no values are given)

Sample code upload TBD

Submission requirements

1. Submit one version of the simulator that **only adds the multiple servers feature**. You must use STL priority queue for event list and STL queue for buffer. Input/Output file can be the same - just skip over the unnecessary inputs (use only the first speed as the speed of the server), and do not print out the speed related metrics
2. Submit a second version which implements everything as specified.
 - a. In this version, along with the code you must submit a presentation (in pdf format) which outlines the class design, the data structures used, and the design decisions taken for efficiency. Detailed instructions regarding this will be posted later.

Evaluation Criteria

This assignment will be judged very strongly for software design. Design & code will be judged for

- Efficiency - in space and time (choice of data structures and algorithms. Other optimizations.).
 - Specific requirement: The implementation of event list should be as a heap with "locator" capability. That is, there must be a way to directly access a specific item and change its priority value (timestamp, in this case). This change may change relative locations of other entries in the heap - your implementation must take care of this issue. "Search" based solutions are semantically incorrect and will face marks penalties.
- Maintainability, extendability, elegance
 - File organization, modularity, encapsulation
 - Will design choices not "scale" with added simulation features?
 - Scalability will be judged both on coding ease vs time/space scalability
- Readability of code
 - Comments, variable names, indentation
- Any other features you wish to "show-off"

Note that no one design is perfect and there are trade-offs. Your presentation should describe the thinking that went into your choice (E.g. "I chose this design because it is efficient, although less maintainable/extendable", or vice-versa)

Some Do's and Don'ts About Simulation

The basic simulation design should always be as follows:

simulation main:

```
initialize simulation {
    set time=0,
    schedule initial events
        (that means create Event objects, timestamp them, add any necessary information
        about which entities in the actual system these events might be about.
        Then throw the event into the event list)
}

while (eventlist not empty) {

    remove the nearest event (pop from priority queue, key being event timestamp, priority
    queue is required because events are not added into the queue in the order that they
    may happen. You may add an event e1 scheduled at time t+s, before you add an event e2
    scheduled at time t.)
    advance simulation clock to this event time
    handle event based on event type (i.e. call event handler).
        (This will usually be a switch that picks the event handler based on the event
        type. For advanced simulations, event handlers will need arguments. This
        could be information that was retrieved from the event object itself. Mainly,
        simulations of complex systems need information about which object of the
        system is the event about. E.g. in multiple server simulation, if you get a
        departure
        event, you need to know which server the departure is from. Or which request
        departing - server ID could be in the request)

}
```

event handlers should have logic like this:

```
update system state - whatever needs to be done to the state, do it.
Based on new system state, new events may get created.
    Create them, add them into eventlist.
Collect statistics.
```

- Do not distribute events in multiple data structures. Consolidate events into **one** event list

- Think about the best way of attaching information to an event, that helps you process it.